

Escriba

Clean, anonymized Markdown for LLMs

Complete manual — users & administrators

Developed by Diego Parrás · CeMIACE · SEUBES · FCE-UBA

github.com/diegoparras/escriba

Part 1 — For users

What is Escriba?

Escriba converts almost any file into clean Markdown — plain text with formatting, ready to copy, save or feed to an AI. It handles documents, spreadsheets, slides, images, audio, video and YouTube.

Getting started

- Enter the password and sign in (your level: HUMANO, ANGEL or DIOS).
- Drag files onto the centre area, or click “Select files”.
- Click “Convert all”. A progress bar appears.
- Open the result and use Preview / Markdown / Split. Copy it or download the .md (or a .zip for several files).

Key features

- Audio & video are transcribed locally with Whisper; pick the language under Advanced options.
- Paste a YouTube link to get its transcript.
- Text in images is OCR'd automatically; scanned PDFs are detected, OCR'd and auto-rotated on the fly.
- AI is optional (OpenAI, Gemini, OpenRouter) and defaults to “No AI”. Models are listed automatically.
- PII anonymization for LLMs: a local engine (NER + invoice layout + validated detectors + your own rules) with five output modes — typed, anonymous, reversible pseudonymization (token → send to the LLM → re-hydrate), partial masking (••••-3456) and stable hashing (same data → same pseudonym across documents).
- Visual redaction: download your PDF or scanned image with the PII blacked out on the page — true redaction: the text and the pixels underneath are removed from the file.
- Queue: convert everything or only the selected files.
- Light/dark theme and an interface available in 7 languages.

Part 2 — Roles & access levels

Login is required. Each level has its own password and limits, all configurable via environment variables.

Capability	HUMANO	ANGEL	DIOS
Upload & convert files	Yes	Yes	Yes
Convert from a public URL	No	Yes	Yes
Internal URL / local path	No	No	Yes
Audio / video / ZIP	No	Yes	Yes
YouTube transcripts	Yes	Yes	Yes
OCR (forced / automatic)	No	Yes	Yes
Use server AI keys	No	Yes	Yes
Max file size	25 MB	100 MB	unlimited
Files per batch	3	10	unlimited
Rate limit (req/min)	15	60	unlimited

Part 3 — For administrators

Architecture

Escriba runs in a single, self-contained Docker container that bundles FastAPI + Markdown, ffmpeg, Tesseract + OCRmyPDF, faster-whisper and an embedded Redis. It runs as a non-root user with a healthcheck and automatic CPU-core detection.

Deployment

Run the prebuilt image (no build needed):

```
docker run -d --name escriba -p 8000:8000 \
  -e SECRET_KEY="<long-random-key>" \
  -e GOD_PASSWORD="<your-password>" \
  ghcr.io/diegoparras/escriba:latest
```

Or with Docker Compose:

```
git clone https://github.com/diegoparras/escriba.git
cd escriba
cp .env.example .env
docker compose up -d --build
```

EasyPanel/Portainer/Dokploy: create an App from the Docker image ghcr.io/diegoparras/escriba:latest, add the environment variables, map the domain to container port 8000 with HTTPS, and deploy.

Environment variables

Variable	Default	Description
SECRET_KEY	random	Session signing key. Set it in production.
GOD/ANGEL/HUMAN_PASSWORD	—	Password for each access level.
HUMAN_OPEN	false	Allow the HUMANO level without login.
WEB_CONCURRENCY	auto	Parallel workers (auto = CPU cores).
MAX_UPLOAD_MB	100	Maximum upload size (except DIOS).
WHISPER_MODEL	base	Transcription model (tiny ... large-v3).
MAX_MEDIA_MINUTES	120	Audio/video length cap (0 = unlimited).
OPENAI/OPENROUTER/GOOGLE_API_KEY	—	Server-side AI keys (fallback).
API_TOKEN / API_TOKEN_ROLE	— / angel	Token for automation and its role.
EMBEDDED_REDIS	true	Embedded Redis for shared rate limiting.
ENABLE_DOCS	false	Expose Swagger at /api/docs.

Server-side AI keys

- If set, users don't need to paste their own key. Only DIOS and ANGEL use them.
- The cost is on you and is shared. Leave them empty so each user brings their own.
- A user-supplied key always takes priority over the server key.

Transcription (Whisper)

- Set the model with WHISPER_MODEL (larger = more accurate, slower).
- It runs on CPU. MAX_MEDIA_MINUTES caps the duration (DIOS has no cap).

API token (automation)

For n8n or scripts, set API_TOKEN and send it as a header:

```
curl -H "X-API-Key: YOUR_TOKEN" \
  -F "file=@document.pdf" \
  https://your-domain/api/convert
```

Security

- Authentication required; three roles with distinct permissions.
- Anti-SSRF (internal IPs/redirects blocked); local paths only for DIOS.
- XSS sanitization (DOMPurify) and security headers (CSP).
- Non-root container with no-new-privileges; OCR/Whisper run without a shell.
- Per-role rate limiting; the API key is never stored on the server.

Privacy

Uploaded files are processed in a temporary file and deleted as soon as the conversion finishes (success or error). There is no database and no uploads folder; the result is only sent to your browser.

Part 4 — API reference

Endpoint	Method	Description
/api/login	POST	Log in; sets a session cookie.
/api/me	GET	Current role and capabilities.
/api/convert	POST	Convert a file or URL to Markdown.
/api/models	POST	List the provider's models.
/api/stats	GET	Server stats (by role).

POST /api/convert (multipart/form-data): file or url, plus optional lang, ocr, llm_provider, llm_api_key, llm_model.

Part 5 — Troubleshooting

- A long video won't transcribe: it may exceed MAX_MEDIA_MINUTES — raise the cap or use DIOS.
- A YouTube video fails: it may have no captions, or YouTube blocks your server IP — paste your YouTube cookies in Settings, or set YT_PROXY/YT_COOKIES.
- Error 413 on upload: raise MAX_UPLOAD_MB and your reverse proxy's body limit.
- A PDF came out empty: it was scanned without OCR for your role — use ANGEL/DIOS or tick "Force OCR".
- Forgot the DIOS password: if unset, a random one is printed in the container logs at startup.

Escriba — by Diego Parrás · CeMIACE · SEUBES · FCE-UBA · MIT License